

Setup Guide — Run the Handbook Generator (PDF → RAG → Chat)

1) Prerequisites

- Python 3.10+ (3.11 recommended)
 - A Supabase project (Postgres)
 - Internet access (for embedding model download on first run)
 - (Optional) xAI API key for Grok (`XAI_API_KEY`)
-

2) Install dependencies

Option A: Use `requirements.txt` (recommended)

Make sure your `requirements.txt` includes at least:

```
streamlit
python-dotenv
supabase
pdfplumber
sentence-transformers
openai
tqdm
numpy
torch>=2.2.0
transformers>=4.43.0
datasets
einops>=0.8.0
```

Then run:

```
pip install -r requirements.txt
```

Option B: Install manually

```
pip install streamlit python-dotenv supabase pdfplumber sentence-transformers openai tqdm  
numpy "torch>=2.2.0" "transformers>=4.43.0" datasets "einops>=0.8.0"
```

3) Create .env file

Create a file named `.env` in the repo root:

```
SUPABASE_URL=YOUR_SUPABASE_URL  
SUPABASE_SERVICE_KEY=YOUR_SUPABASE_SERVICE_KEY  
  
# Optional: Grok via xAI (OpenAI-compatible)  
XAI_API_KEY=YOUR_XAI_API_KEY  
# or GROK_API_KEY=YOUR_XAI_API_KEY  
  
GROK_MODEL=grok-4-1-fast-reasoning  
GROK_MAX_TOKENS=4000
```

Where to find Supabase values:

- Supabase → **Project Settings** → **API**
 - Project URL → `SUPABASE_URL`
 - service_role key → `SUPABASE_SERVICE_KEY`
-  Keep the service key private (never expose client-side).
-

4) Supabase setup (Required)

Open Supabase → **SQL Editor** and run these in order.

4.1 Enable pgvector

create extension if not exists vector;

4.2 Create tables

```
create table if not exists public.documents (  
    id uuid primary key default gen_random_uuid(),  
    filename text not null,  
    created_at timestampz default now()
```

```

);

create table if not exists public.chunks (
    id uuid primary key default gen_random_uuid(),
    document_id uuid references public.documents(id) on delete cascade,
    chunk_index integer not null,
    content text not null,
    metadata jsonb default '{}':jsonb,
    embedding vector(384) not null
);

create index if not exists chunks_document_id_idx on public.chunks(document_id);

```

4.3 Create retrieval function (RPC)

```

create or replace function public.match_chunks(
    query_embedding vector,
    match_count integer,
    filter_doc uuid default null
)
returns table (
    id uuid,
    document_id uuid,
    chunk_index integer,
    content text,
    metadata jsonb,
    similarity double precision
)
language sql
stable
as $$
select
    c.id,
    c.document_id,
    c.chunk_index,
    c.content,
    c.metadata,
    1 - (c.embedding <=> query_embedding) as similarity
from public.chunks c
where (filter_doc is null or c.document_id = filter_doc)
order by c.embedding <=> query_embedding
limit match_count;
$$;

```

5) Run the app

From the repo root:

```
streamlit run app/main.py
```

It will open in your browser.

6) How to use

6.1 Index a PDF

1. Upload a **text-based PDF** in the sidebar
2. Click **Index PDF**
3. You should see: **Indexed ✓ document_id=...**

6.2 Ask questions (RAG chat)

Ask something in the PDF:

- You should get an answer with citations like **(PDF p. 2)**
Ask something not in the PDF:
- You should get: **"The uploaded PDFs don't mention this."**

6.3 Generate a handbook

In chat:

```
/handbook Retrieval-Augmented Generation
```

Download the output from the sidebar as Markdown.

7) Troubleshooting

PowerShell blocks venv activation

Run:

```
Set-ExecutionPolicy -ExecutionPolicy RemoteSigned -Scope CurrentUser
```

Or activate via cmd:

```
venv\Scripts\activate.bat
```

Retrieval returns no results

Confirm chunks exist:

```
select count(*) from chunks;
```

Confirm RPC works:

```
select * from match_chunks(  
    (select embedding from chunks limit 1),  
    5,  
    (select document_id from chunks limit 1)  
)
```

First run is slow

The embedding model downloads on the first run; later runs are much faster.